




# Machine Learning 5: Convolutional Neural Networks



# Recap




# Recap




- A neural network is made up of a series of layers with artificially created 'features' or **activations**.

# Recap




- A neural network is made up of a series of layers with artificially created 'features' or **activations**.
- These features are created through **non-linearly transforming** (e.g. sigmoid function) a bf weighted combination of the input features.

# Recap



- A neural network is made up of a series of layers with artificially created 'features' or **activations**.
- These features are created through **non-linearly transforming** (e.g. sigmoid function) a bf weighted combination of the input features.
- The **training process** consists of a series of **forward propagations** (to determine the cost) and **back propagations** (to determine the cost derivatives) in order to optimise the weights.

# Recap



- A neural network is made up of a series of layers with artificially created 'features' or **activations**.
- These features are created through **non-linearly transforming** (e.g. sigmoid function) a bf weighted combination of the input features.
- The **training process** consists of a series of **forward propagations** (to determine the cost) and **back propagations** (to determine the cost derivatives) in order to optimise the weights.
- Typically each hidden layer of a neural network is **fully connected**, i.e. each neuron in layer  $N$  is connected to each neuron in layer  $N \pm 1$ .

# A short-coming of standard neural networks

In the context of image classification, a neural network with fully connected layers can end up over-fitting very easily.

# A short-coming of standard neural networks

In the context of image classification, a neural network with fully connected layers can end up over-fitting very easily.


For example a  $200 \times 200$  pixed color (RGB:  $\times 3$ ) image produces 120,000 weights for each neuron of our first hidden layer.



# A short-coming of standard neural networks

In the context of image classification, a neural network with fully connected layers can end up over-fitting very easily.

For example a  $200 \times 200$  pixel color (RGB:  $\times 3$ ) image produces 120,000 weights for each neuron of our first hidden layer.



# Convolutions


One way to solve this problem is to make use of the fact that our input is an image. This image can be represented as a volume, for example a  $200 \times 200 \times 3$  cuboid.

Segments of this volume can then be convolved with weight arrays to produce a new volume. We can perform successive such convolutions to finally ending up with a reduced output volume, for example  $1 \times 2$  for the face-not face classification.

# Convolutions

One way to solve this problem is to make use of the fact that our input is an image. This image can be represented as a volume, for example a  $200 \times 200 \times 3$  cuboid.

Segments of this volume can then be convolved with weight arrays to produce a new volume. We can perform successive such convolutions to finally ending up with a reduced output volume, for example  $1 \times 2$  for the face-not face classification.



**Credit:** <https://cs231n.github.io/convolutional-networks/>

# How does this work in practice?

# How does this work in practice?

Each weight array, or **filter**, scans the image jumping  $N$ -pixels at a time ( $N=1$  usually).  $N$  is called the stride.

# How does this work in practice?

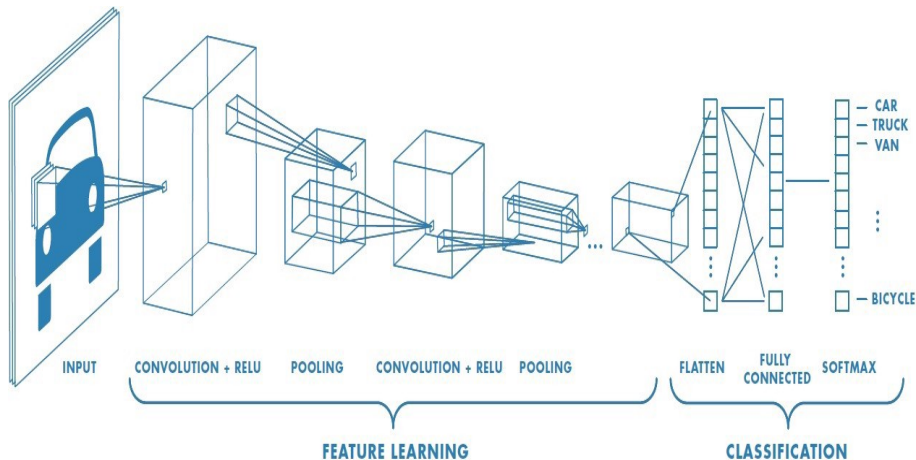
Each weight array, or **filter**, scans the image jumping  $N$ -pixels at a time ( $N=1$  usually).  $N$  is called the stride.

When it is applied to a segment of the image an element-wise multiplication is performed and then the products are summed to give the output 'pixel'.

**See animation at at**

<https://cs231n.github.io/convolutional-networks/>

# The layers of a CNN



Credit: [https://medium.com/@\\_sumitsaha\\_](https://medium.com/@_sumitsaha_)

# The layers of a CNN

**Conv layer:** The filters to be applied to the input volume with a depth equal to the input volume depth (in our example 3 - RGB). These filters are applied to segments of the input volume sequentially, producing a new volume.



# The layers of a CNN

**Conv layer:** The filters to be applied to the input volume with a depth equal to the input volume depth (in our example 3 - RGB). These filters are applied to segments of the input volume sequentially, producing a new volume.

**Pooling layer:** A pooling layer acts to reduce the volume and hence number of weights for a successive Conv layer. It does this by moving through the input volume at a stride greater than 1 pixel at a time. One example is max pooling which applies a Max function to select most important input features.

**Flatten and Dense layer:** Unrolls the volume into a vector of values and then the dense layer is a fully connected normal neural network hidden layer with neurons equal to the number of classes.

**Flatten and Dense layer:** Unrolls the volume into a vector of values and then the dense layer is a fully connected normal neural network hidden layer with neurons equal to the number of classes.

**Softmax:** Apply softmax function to generate probabilities for classes.

$$\sigma_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}. \quad (1)$$

**Flatten and Dense layer:** Unrolls the volume into a vector of values and then the dense layer is a fully connected normal neural network hidden layer with neurons equal to the number of classes.

**Softmax:** Apply softmax function to generate probabilities for classes.

$$\sigma_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}. \quad (1)$$

**Dropout:** Dropout is a technique which acts as a regularisation of the network. During training, random activations are not considered in the layer prior to a dropout operation. This creates 'noise' in the training process. By forcing the weights to account for this, we make each individual weight less important, and consequently we reduce over-fitting.

**Flatten and Dense layer:** Unrolls the volume into a vector of values and then the dense layer is a fully connected normal neural network hidden layer with neurons equal to the number of classes.

**Softmax:** Apply softmax function to generate probabilities for classes.

$$\sigma_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_{j=1}^{j=K} e^{z_j}}. \quad (1)$$

**Dropout:** Dropout is a technique which acts as a regularisation of the network. During training, random activations are not considered in the layer prior to a dropout operation. This creates 'noise' in the training process. By forcing the weights to account for this, we make each individual weight less important, and consequently we reduce over-fitting.

You can now have a look at 'tutorials/NN\_tutorial.ipynb' and [BaCoN!](#)